



#ASLI ENGINEERING

GitHub Outage Downtime due to an Edge Case



BY

ARPIT BHAYANI

Dissecting GitHub Outage

Edge case and a poor SQL query

August 2021. GitHub experienced major outage that affected their core services

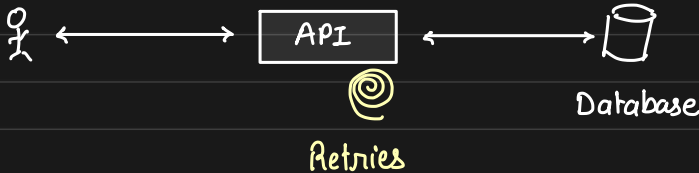
What happened?

The outage was due to an edge case that generated poor SQL query



One of their MySQL masters degraded

+ automatic retry and queueing made it worse



We add retries in order to cover transient issues but it can also become something that adds to load during an outage

So, how could an "edge case" "degrade" the master?


Note: This is just an example, nothing to do with GitHub

Just trying to explain how an edge case can put huge load on DB

Say, we have **commits** table

commits
id, msg, author_id, created_at

epoch ms
stored as integer



```
SELECT count(id) FROM commits  
WHERE author_id = 123  
AND created_at > start_date
```

Say, the service is written in Golang

and we get **start_date** in the query parameter
↓
as epoch seconds

So, from the query parameter, we have to
convert it to integer and then pass it to
our SQL query

`https://api.github.com/stats/commits?start_date=<epoch-ts>`

Golang based code to do this could look something like this

```
int start_date;  
  
start_date, __ := strconv.Atoi(req.q.start_date)  
  
res = query.execute(  
    "SELECT count(id) FROM commits  
    WHERE author_id = 123  
    AND created_at > " + start_date)
```

Edge case: what if start_date is not passed

what if non-integer value is passed as start_date

↓

Some misconfiguration on "calendar" widget

`strconv.Atoi()` → throws an error

but we did not handle it

So, 'start_date' being integer will get the

default value which is 0.

the default value is set `int start-date;`
 ↓
 default value = 0

The query that gets fired is

```
SELECT count(id) FROM commik  
WHERE author-id = 123  
AND created_at > 0;
```

This would literally count **all the commits** made by the user ... ever!
and would take a long of time to execute
and put some load on DB

Given this is a common query, it will be invoked frequently
and will put massive load on DB!!
causing degradation and eventually an outage

So, how to protect ourselves from such situation?

So, how to avoid such situations

→ Be pessimistic

1. Sanitize your input before query execution
2. Apply hard guard rails. even in edge case it goes through
say 1000 rows

```
SELECT count(id) FROM commik  
WHERE author-id = 123  
AND created-at > 0  
LIMIT 1000;
```