



**#ASLI ENGINEERING**

# Realtime Data Aggregation with DynamoDB



**BY**

**ARPIT BHAYANI**

## Realtime Data Aggregation

with DynamoDB

## What is DynamoDB?

DynamoDB is low-latency KV store offered by AWS.



Operations it supports are GET, SET, PARTIAL UPDATE, DEL, TTL

## Deliveroo Data Model

Deliveroo does food delivery and it is quite popular in Europe.

1. Restaurant owners list their restaurants
2. Users can browse through them, pick their favourite, and order


Users can also mark some restaurants as **FAVOURITE**.

## favowrites table

restaurant\_id \_\_ user\_id : Partition key

This restaurant 'favourite' of this user

created-at } other attributes  
user-id }



1. Getting if a user 'u' has restaurant 'r' in favourite


Is just a simple GET (kv lookup)

2. Marking a favourite or unmarking are pointed operations

Thus we see, all user operations happen in O(1) time

Feature request: Show list of restaurants ordered by FAV counts

Hence we would need to "aggregate" favourites and keep the favourite count handy at a restaurant level.



Core Idea: Consume all the updates happening in Favourites table and aggregate them in a different table

## Aggregated Favournites

restaurant\_id : partition key

time\_window : sort key

↳ ALL\_TIME , AUG 2022 , 20220801

favournite\_count

updated\_at

A	ALL_TIME	729
A	AUG 2022	42
B	ALL_TIME	10072

## Aggregator function would

- receive a batch of updates happening in Favournites table
- filter INSERT and REMOVE events
- **atomically**, increment or decrement the count per restaurant

↳ use ADD instead of SET

↳ use DB transactions to batch update 100 writes

Isn't lambda costly?

Not a frequent action at all

# Fav per day ~ 7000

Each execution ~ 25ms

Mem consumption 64MB

< \$1 per month

## Advantages

- extremely cost efficient for given user behaviour
- APIs response time unaffected because ASYNC
- better than traditional approaches
  - daily full table scan (not realtime)
  - synchronous update to Aggregated Favourite table
    - ↳ (APIs would become slower)