



**#ASLI ENGINEERING**

# How Twitter scales and keeps their search stable



**BY**

**ARPIT BHAYANI**

## Scalability and Stability of Search @ Twitter

Twitter uses Elasticsearch to power search of tweets, user, direct messages




Why Elasticsearch?

Speed, Scalability, distributed nature, and simple REST APIs

Internal customers / services directly talk to ES and get things done

Hence, a need of standardization, stability and performance.



Elasticsearch Proxy

Querying, indexing, monitoring, metrics was all manual & to be done separately. Proxy standardized throttling, routing, security, authentication.


monitoring - cluster health, success rate, failure rate, latency



## Ingestion Service

When there is a massive surge in traffic,

Elasticsearch gives up!




↳ increased indexing latency

↳ increased query latency

To handle heavy ingestion gracefully, Twitter built an ingestion service

Ingestion service queues indexing request in kafka

later to be consumed by workers



## Advantages :

**Request batching** : batch write on ES cluster

**Backpressure** : consuming at its own pace

**Throttling** : slowing down if ES is overwhelmed

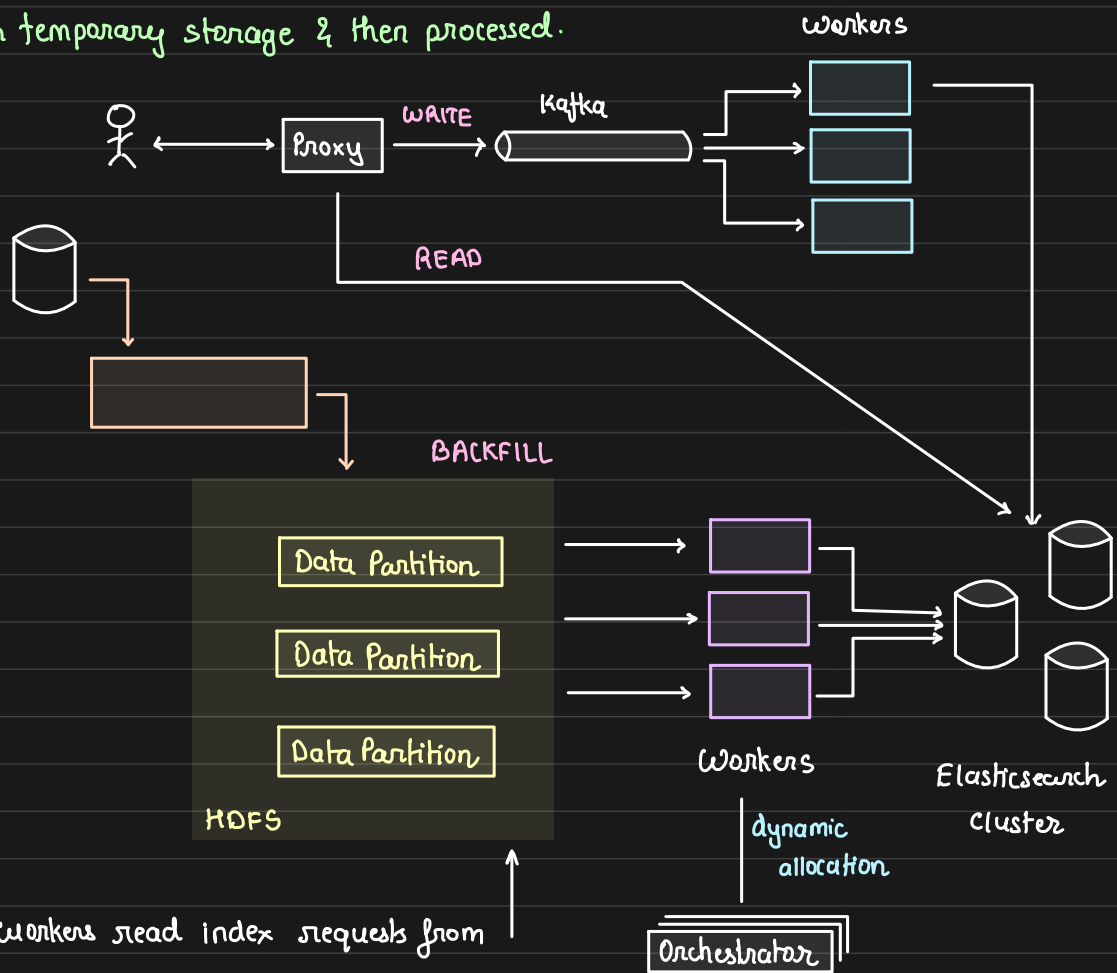
**Retries** : if cluster down, we can retry easily

Backfill Service (ingesting 100s of TB of data in ES)

Directly ingesting massive data using MapReduce synchronously does not work

Indexing happened in sync from MapRed. Elasticsearch cannot handle this huge ingestion in one shot.

Now the indexing request is dumped in temporary storage & then processed.



Workers read index requests from and index data into the cluster.